

LECTURE 4

Unit 2

WHAT IS SQL?

- ◉ **SQL - Structured Query Language**
- ◉ Developed by IBM
- ◉ It is a common database language which has statements for both data definitions, query, and update
- ◉ It is both a DDL (**Data Definition Language**) and DML (**Data Manipulation Language**)
- ◉ SQL uses the terms table, row, and column for the formal relational model terms *relation*, *tuple*, and *attribute*, respectively

PROS AND CONS

⦿ Pros

- Very flexible
- Universal (Oracle, Access, etc)
- Few Commands to Learn

⦿ Cons

- Requires Detailed Knowledge of the Structure of the Database
- Can Provide Misleading Results

CHARACTERISTICS OF SQL

- ◉ SQL allows the user to create, update, delete, and retrieve data from a database.
- ◉ SQL is very simple and easy to learn.
- ◉ SQL works with database programs like DB2, Oracle, MS Access, MS SQL Sever etc.

GENERAL DATA TYPES IN SQL

- ◉ CHARACTER (n) - character string of fixed length n
- ◉ VARCHAR (n) - character string of variable length, maximum length is n
- ◉ BOOLEAN - stores TRUE or FALSE values
- ◉ INTEGER and SMALLINT - integer number
- ◉ FLOAT and REAL - floating point number
- ◉ BIT (n) - array of n bits
- ◉ DATE - stores year, month and day values
- ◉ TIME - stores hour, minute and second values
- ◉ TIMESTAMP - stores year, month, day, hour, minute and second values

TYPES OF SQL COMMANDS

- SQL commands are instructions used to communicate with the database to perform specific task that work with data.
- Two common SQL commands
 - DDL
 - DML

DATA DEFINITION LANGUAGE (DDL)

- ⦿ These SQL commands are used for creating, modifying, and dropping the structure of database objects.
- ⦿ The most basic items of DDL are:
 - **CREATE** creates an object (a table, for example) in the database
 - **DROP** deletes an object in the database
 - **ALTER** modifies the structure an existing object in various ways—for example, adding a column to an existing table.

DATA MANIPULATION LANGUAGE (DML)

- ⦿ These SQL commands are used for storing, retrieving, modifying, and deleting data.
- ⦿ These commands are
 - SELECT
 - INSERT
 - UPDATE
 - DELETE

CREATE SCHEMA

- ◉ It defines a database schema

- ◉ Syntax

**CREATE SCHEMA schema-name AUTHORIZATION
user-name**

- ◉ Example

CREATE SCHEMA LIBRARY;

**CREATE SCHEMA LIBRARY AUTHORIZATION
James;**

CREATE TABLE

- ◉ Create a table in the database

- ◉ Syntax

`CREATE TABLE table-name (col-name1 data-type,
co-name2 data-type,...);`

- ◉ Example

**`CREATE TABLE Student (Name VARCHAR(40),
RollNo INTEGER , Class INTEGER);`**

SQL CONSTRAINTS

- ⦿ Used to limit the type of data that can go in a table
- ⦿ Important ones
 - PRIMARY KEY
 - NOT NULL
 - UNIQUE
 - DEFAULT
 - FOREIGN KEY

◎ PRIMARY KEY

- This constraint uniquely identifies each record (row) in a table

```
CREATE TABLE STUDENT (Name VARCHAR(40),  
RollNo INTEGER PRIMARY KEY, ClassNo INTEGER);
```

◉ FOREIGN KEY

- It points to a primary key in another table

```
CREATE TABLE STUDENT (Name VARCHAR(40),  
RollNo INTEGER, ClassNo INTEGER FOREIGN KEY  
REFERENCES CLASS (ClassNo));
```

⦿ NOT NULL

- Enforces a column to not accept NULL values

```
CREATE TABLE STUDENT (Name VARCHAR(40) NOT  
NULL, RollNo INTEGER, ClassNo INTEGER);
```

◎ UNIQUE

- Similar to PRIMARY KEY
- Uniquely identifies each record in a table

```
CREATE TABLE STUDENT (Name VARCHAR(40),  
RollNo INTEGER UNIQUE, ClassNo INTEGER);
```

◎ DEFAULT

- Used to insert a default value in a column

```
CREATE TABLE STUDENT (Name VARCHAR(40)  
DEFAULT 'James', RollNo INTEGER, ClassNo  
INTEGER);
```


DROP TABLE

- ◉ Remove / delete a table from a database along with its definition

DROP TABLE EMPLOYEE;

ALTER TABLE

- ◉ Add, delete or modify columns in an existing table

- ◉ To add a column

```
ALTER TABLE STUDENT ADD Address VARCHAR(40);
```

- ◉ To delete a column

```
ALTER TABLE STUDENT DROP COLUMN Address;
```

- To change a data type of a column

```
ALTER TABLE STUDENT ALTER COLUMN RollNo  
CHAR(10);
```

INSERT INTO

- ◉ Insert a new row in a table

- ◉ Two forms

- Do not specify column names

```
INSERT INTO STUDENT VALUES ('John', 12, 4);
```

- Specify column names and values

```
INSERT INTO STUDENT (Name, RollNo, ClassNo)  
VALUES ('John', 12, 4);
```

UPDATE

- Update existing records in a table

- Syntax

UPDATE table-name

SET col1=value, col2=value,...

WHERE some-col=some-value;

◉ Example

```
UPDATE STUDENT  
SET Name='Mary'  
WHERE RollNo=3;
```

DELETE

- ◉ Delete rows in a table

- ◉ Syntax

```
DELETE FROM table-name  
WHERE some-col=some-val;
```

- ◉ Example

```
DELETE FROM STUDENT  
WHERE ClassNo=4;
```

SELECT

- ◉ Used for retrieving rows from database
- ◉ General syntax
 - SELECT column-name-list
 - FROM table-name-list
 - WHERE condition;
- ◉ SELECT - attributes to extract
- ◉ FROM - tables from where to extract
- ◉ WHERE - optional, specifies the condition(s)

- Retrieve names and roll numbers of students

```
SELECT Name, RollNo  
FROM STUDENT;
```

- Retrieve all details of students

```
SELECT *  
FROM STUDENT;
```

- Retrieve roll numbers of all students belonging to class 6.

```
SELECT RollNo  
FROM STUDENT  
WHERE ClassNo=6;
```

USING DISTINCT

- ◉ Select only distinct values from a table

```
SELECT DISTINCT ClassNo  
FROM STUDENT;
```

- This statement will list only distinct values of Class Number, even though multiple occurrences may exist

ALIAS

- Alias names can also be given to tables

```
SELECT S.Name, S.RollNo  
FROM STUDENT AS S;
```

SET OPERATIONS

- There is a union operation (**UNION**), and in *some versions* of SQL there are set difference (**MINUS**) and intersection (**INTERSECT**) operations
- The resulting relations of these set operations are sets of tuples; *duplicate tuples are eliminated from the result*
- The set operations apply only to *union compatible relations* ; the two relations must have the same attributes and the attributes must appear in the same order

- Make a list of all project numbers for projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.

```
(SELECT PNAME
FROM   PROJECT, DEPARTMENT, EMPLOYEE
WHERE  DNUM=DNUMBER AND MGRSSN=SSN AND
        LNAME='Smith')

UNION

(SELECT PNAME
FROM   PROJECT, WORKS_ON, EMPLOYEE
WHERE  PNUMBER=PNO AND ESSN=SSN AND
        LNAME='Smith');
```

NESTING OF QUERIES

- ◉ A complete SELECT query, called a *nested query* , can be specified within the WHERE-clause of another query, called the *outer query*

Retrieve the name and address of all employees who work for the 'Research' department.

```
SELECT FNAME, LNAME, ADDRESS
FROM   EMPLOYEE
WHERE  DNO IN
      (SELECT DNUMBER
       FROM   DEPARTMENT
       WHERE  DNAME='Research' );
```

CORRELATED NESTED QUERIES

- ◉ If a condition in the WHERE-clause of a *nested query* references an attribute of a relation declared in the *outer query* , the two queries are said to be *correlated*
- ◉ The result of a correlated nested query is *different for each tuple (or combination of tuples) of the relation(s) the outer query*

Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
SELECT      E.FNAME, E.LNAME
FROM        EMPLOYEE AS E
WHERE E.SSN IN (SELECT ESSN
                FROM   DEPENDENT
                WHERE  ESSN=E.SSN AND
                       E.FNAME=DEPENDENT_NAME);
```


- ⦿ A query written with nested SELECT... FROM... WHERE... blocks and using the = or IN comparison operators can *always* be expressed as a single block query.

```
SELECT      E.FNAME, E.LNAME  
FROM        EMPLOYEE E, DEPENDENT D  
WHERE       E.SSN=D.ESSN AND  
            E.FNAME=D.DEPENDENT_NAME;
```

EXISTS

- EXISTS is used to check whether the result of a correlated nested query is empty (contains no tuples) or not

Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
SELECT      FNAME, LNAME
FROM        EMPLOYEE
WHERE       EXISTS
            (SELECT      *
FROM          DEPENDENT
WHERE         SSN=ESSN AND
FNAME=DEPENDENT_NAME);
```

Retrieve the names of employees who have no dependents.

```
SELECT      FNAME, LNAME  
FROM        EMPLOYEE  
WHERE       NOT EXISTS  
            (SELECT      *  
              FROM        DEPENDENT  
              WHERE SSN=ESSN);
```

ORDER BY

- Used when we want data to appear in a specific order
 - Retrieve all details of students and then arrange the result on the basis of ascending order of class number

```
SELECT *  
FROM STUDENT  
ORDER BY ClassNo;
```

```
SELECT *  
FROM STUDENT  
ORDER BY ClassNo ASC;
```

- Retrieve all details of students and then arrange the result on the basis of descending order of class number

```
SELECT *  
FROM STUDENT  
ORDER BY ClassNo DESC;
```

COMMON OPERATORS

=	Equal
<>	Not equal
>	Greater than
<	Less than
>=	Greater than or equal
<=	Less than or equal
BETWEEN	Between an inclusive range
AND	Display if both conditions are true
OR	Display if either conditions are
true	

- Retrieve all names of students from class 4 to class 6

```
SELECT Name  
FROM STUDENT  
WHERE ClassNo BETWEEN 4 AND 6;
```

JOIN

- Used to query data from two or more tables

Retrieve the student names, roll numbers and class who belong to a class whose class teacher is 'Mary'

```
SELECT S.Name, S.RollNo, C.ClassNo  
FROM STUDENT AS S, CLASS AS C  
WHERE S.ClassNo = C.ClassNo AND C.Teacher =  
'Mary';
```


SAMPLE TABLE

EmployeeStatisticsTable			
EmployeeIDNo	Salary	Benefits	Position
010	75000	15000	Manager
105	65000	15000	Manager
152	60000	15000	Manager
215	60000	12500	Manager
244	50000	12000	Staff
300	45000	10000	Staff
335	40000	10000	Staff
400	32000	7500	Entry-Level
441	28000	7500	Entry-Level

- *Retrieve the employee numbers of those employees who have salary more than Rs. 50,000*

```
SELECT EmployeeIDNo  
FROM EMPLOYEESTATISTICSTABLE  
WHERE Salary > 50000;
```

- Try the queries below
 - Retrieve employee numbers who has position as 'Manager'
 - Retrieve employee numbers and benefits who has position as 'Manager' and whose salary is above Rs. 50,000

TAKE THESE THREE TABLES AS SHOWN

AntiqueOwners

OwnerID	OwnerLastName	OwnerFirstName
01	Jones	Bill
02	Smith	Bob
15	Lawson	Patricia
21	Akins	Jane
50	Fowler	Sam

Orders

OwnerID	ItemDesired
02	Table
02	Desk
21	Chair
15	Mirror

- Here the SellerID and BuyerID also means they are OwnerID

Antiques		
SellerID	BuyerID	Item
01	50	Bed
02	15	Table
15	02	Chair
21	50	Mirror
50	01	Desk
01	21	Cabinet
02	21	Coffee Table
15	50	Chair
01	15	Jewelry Box
02	21	Pottery
21	02	Bookcase
50	01	Plant Stand

QUERIES

- List all the details of all the Antique owners

SELECT *

FROM ANTIQUEOWNERS;

- Try the queries
 - List all details of all the Orders
 - List all details of all Antiques

- Retrieve all items ordered by the owner whose name is 'Bob Smith'

```
SELECT O.ItemDesired  
FROM ORDERS AS O, ANTIQUEOWNERS AS A  
WHERE A.OwnerFirstName = 'Bob' AND  
A.OwnerLastName = "Smith" AND A.OwnerID  
= O.OwnerID;
```

- Retrieve the names of the buyers who bought the item 'Table'

```
SELECT A.OwnerFirstName, A.OwnerLastname  
FROM ANTIQUEOWNERS AS A, ANTIQUES AS Q  
WHERE Q.Item='Table' AND Q.BuyerID=A.OwnerID;
```

○ Try these queries

- Retrieve the names of sellers who sell 'Mirror' or 'Desk'
- Retrieve the seller IDs of who sell items to buyer ID 50
- Retrieve the seller names who sell items to 'Bill Jones'
- Retrieve the buyer names who buy item from 'Jane Akins'

AGGREGATE FUNCTIONS

- Some of the common functions are
 - SUM() : total of all the rows satisfying a condition (s) of the given column, where the column is numeric
 - MIN() : gives the smallest value in the given column
 - MAX() : gives the highest value in the given column
 - COUNT(*) : gives the number of rows satisfying the condition
 - AVG() : gives the average of the given column

SAMPLE TABLE

EmployeeStatisticsTable

EmployeeIDNo	Salary	Benefits	Position
010	75000	15000	Manager
105	65000	15000	Manager
152	60000	15000	Manager
215	60000	12500	Manager
244	50000	12000	Staff
300	45000	10000	Staff
335	40000	10000	Staff
400	32000	7500	Entry-Level
441	28000	7500	Entry-Level

- Retrieve the total of all salaries of employees
SELECT SUM (Salary)
FROM EMPLOYEEESTATISTICSTABLE;
- Retrieve the average salary of all employees
SELECT AVG (Salary)
FROM EMPLOYEEESTATISTICSTABLE;

- Retrieve the maximum and the minimum benefits of all employees

```
SELECT MIN (Benefits), MAX (Benefits)  
FROM EMPLOYEESTATISTICSTABLE;
```

- List the total number of employees

```
SELECT COUNT (*)  
FROM EMPLOYEESTATISTICSTABLE;
```

- ◉ List the total number of employees who has 'Staff' position

```
SELECT COUNT (*)  
FROM EMPLOYEESTATISTICSTABLE  
WHERE Position = 'Staff';
```

- ◉ List the highest salary of the employee who has position 'Manager'

NESTING OF QUERIES

- ◉ A complete SELECT query, called a *nested query* , can be specified within the WHERE-clause of another query, called the *outer query*
- ◉ Many of the previous queries can be specified in an alternative form using nesting
- ◉ Query 1: Retrieve the name and address of all employees who work for the 'Research' department.

```
Q1:  SELECT FNAME, LNAME, ADDRESS
      FROM  EMPLOYEE
      WHERE DNO IN
      (SELECT DNUMBER
       FROM  DEPARTMENT
       WHERE DNAME='Research' );
```

CORRELATED NESTED QUERIES

- ◉ If a condition in the WHERE-clause of a *nested query* references an attribute of a relation declared in the *outer query* , the two queries are said to be *correlated*
- ◉ The result of a correlated nested query is *different for each tuple (or combination of tuples) of the relation(s) the outer query*
- ◉ Query 12: Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
Q12: SELECT      E.FNAME, E.LNAME
      FROM        EMPLOYEE AS E
      WHERE E.SSN IN (SELECT      ESSN
                      FROM    DEPENDENT
                      WHERE ESSN=E.SSN AND
                            E.FNAME=DEPENDENT_NAME);
```

THE EXISTS FUNCTION

- EXISTS is used to check whether the result of a correlated nested query is empty (contains no tuples) or not

THE EXISTS FUNCTION (CONT.)

- Retrieve the name of each employee who has a dependent with the same first name as the employee.

Q12:

```
SELECT      FNAME, LNAME
FROM EMPLOYEE
WHERE       EXISTS
(SELECT     *
FROM DEPENDENT
WHERE      SSN=ESSN AND
FNAME=DEPENDENT_NAME);
```

THE EXISTS FUNCTION (CONT.)

- ◉ Query 6: Retrieve the names of employees who have no dependents.

Q6:

```
SELECT      FNAME, LNAME
FROM EMPLOYEE
WHERE NOT EXISTS
(SELECT      *
FROM DEPENDENT
WHERE SSN=ESSN);
```

- In Q6, the correlated nested query retrieves all DEPENDENT tuples related to an EMPLOYEE tuple. If *none exist*, the EMPLOYEE tuple is selected
- EXISTS is necessary for the expressive power of SQL

EXPLICIT SETS

- It is also possible to use an **explicit (enumerated) set of values** in the WHERE-clause rather than a nested query
- Query 13: Retrieve the social security numbers of all employees who work on project number 1, 2, or 3.

```
Q13:      SELECT      DISTINCT ESSN
           FROM WORKS_ON
           WHERE        PNO IN  (1, 2, 3);
```

NULLS IN SQL QUERIES

- SQL allows queries that check if a value is NULL (missing or undefined or not applicable)
- SQL uses **IS** or **IS NOT** to compare NULLs because it considers each NULL value distinct from other NULL values, so equality comparison is not appropriate .
- Query 14: Retrieve the names of all employees who do not have supervisors.

```
Q14:SELECT      FNAME, LNAME  
      FROM      EMPLOYEE  
      WHERE     SUPERSSN IS NULL;
```

- Note: If a join condition is specified, tuples with NULL values for the join attributes are not included in the result

GROUPING

- ◉ In many cases, we want to apply the aggregate functions *to subgroups of tuples in a relation*
- ◉ Each subgroup of tuples consists of the set of tuples that have *the same value* for the *grouping attribute(s)*
- ◉ The function is applied to each subgroup independently
- ◉ SQL has a **GROUP BY**-clause for specifying the grouping attributes, which *must also appear in the SELECT-clause*

GROUPING (CONT.)

- ◉ Query 20: For each department, retrieve the department number, the number of employees in the department, and their average salary.

```
Q20: SELECT      DNO, COUNT (*), AVG (SALARY)
      FROM  EMPLOYEE
      GROUP BY    DNO;
```

GROUPING (CONT.)

- ◉ Query 21: For each project, retrieve the project number, project name, and the number of employees who work on that project.

```
Q21: SELECT PNUMBER, PNAME, COUNT (*)  
      FROM PROJECT, WORKS_ON  
      WHERE PNUMBER=PNO  
      GROUP BY PNUMBER, PNAME;
```

THE HAVING-CLAUSE

- ◉ Sometimes we want to retrieve the values of these functions for only those *groups that satisfy certain conditions*
- ◉ The HAVING-clause is used for specifying a selection condition on groups (rather than on individual tuples)

THE HAVING-CLAUSE (CONT.)

- ◉ Query 22: For each project *on which more than two employees work* , retrieve the project number, project name, and the number of employees who work on that project.

Q22: SELECT PNUMBER, PNAME, COUNT (*)
 FROM PROJECT, WORKS_ON
 WHERE PNUMBER=PNO
 GROUP BY PNUMBER, PNAME
 HAVING COUNT (*) > 2;

KINDLY GO THROUGH
THESE TOPICS

SUBSTRING COMPARISON

- ◉ The **LIKE** comparison operator is used to compare partial strings
- ◉ Two reserved characters are used: '%' (or '*' in some implementations) replaces an arbitrary number of characters, and '_' replaces a single arbitrary character

SUBSTRING COMPARISON (CONT.)

- ◉ Query 25: Retrieve all employees whose address is in Houston, Texas. Here, the value of the ADDRESS attribute must contain the substring 'Houston,TX'.

Q25:	SELECT	FNAME, LNAME
	FROM	EMPLOYEE
	WHERE	ADDRESS LIKE
		'%Houston,TX%';

SUBSTRING COMPARISON (CONT.)

- ◉ Query 26: Retrieve all employees who were born during the 1950s. Here, '5' must be the 8th character of the string (according to our format for date), so the BDATE value is '_____5_', with each underscore as a place holder for a single arbitrary character.

```
Q26:SELECT      FNAME, LNAME
      FROM EMPLOYEE
      WHERE      BDATE LIKE
      '_____5_';
```

- ◉ The LIKE operator allows us to get around the fact that each value is considered atomic and indivisible; hence, in SQL, character string attribute values are not atomic

ARITHMETIC OPERATIONS (

- The standard arithmetic operators '+', '-', '*', and '/' (for addition, subtraction, multiplication, and division, respectively) can be applied to numeric values in an SQL query result
- Query 27: Show the effect of giving all employees who work on the 'ProductX' project a 10% raise.

```
Q27: SELECT      FNAME, LNAME, 1.1*SALARY
                FROM  EMPLOYEE, WORKS_ON,
PROJECT
                WHERE SSN=ESSN AND PNO=PNUMBER AND
                   PNAME='ProductX';
```

SUMMARY OF SQL QUERIES

- A query in SQL can consist of up to six clauses, but only the first two, SELECT and FROM, are mandatory. The clauses are specified in the following order:

SELECT <attribute list>
FROM <table list>
[WHERE <condition>]
[GROUP BY <grouping attribute(s)>]
[HAVING <group condition>]
[ORDER BY <attribute list>]

SUMMARY OF SQL QUERIES (CONT.)

- ◉ The SELECT-clause lists the attributes or functions to be retrieved
- ◉ The FROM-clause specifies all relations (or aliases) needed in the query but not those needed in nested queries
- ◉ The WHERE-clause specifies the conditions for selection and join of tuples from the relations specified in the FROM-clause
- ◉ GROUP BY specifies grouping attributes
- ◉ HAVING specifies a condition for selection of groups
- ◉ ORDER BY specifies an order for displaying the result of a query
- ◉ A query is evaluated by first applying the WHERE-clause, then GROUP BY and HAVING, and finally the SELECT-clause